

Gazi Aman Khan

AI-ML Engineer | Data Science Associate

✉ gazikhaan15@gmail.com ☎ 9521857892 📍 Udaipur 🌐 [Gazi Aman Khan](#) 🔄 [Gazi Aman Khan](#)

PROFILE

AI/ML Engineer with hands-on experience in Python, SQL, Machine Learning, and Generative AI. Built predictive models, NLP solutions, and a deployed RAG-based application using FastAPI, Gemini, Qdrant, and Streamlit. Passionate about leveraging data and AI to solve real-world problems through scalable, user-focused solutions.

EDUCATION

B.Tech in Computer Science , Techno India NJR Institute of Technology	2021 – 2025 Udaipur
Senior Secondary(XII), CBSE , Central Academy School	2019 – 2020 Udaipur

SKILLS

Languages & Tool — Python (Scikit-learn, Pandas, NumPy, Matplotlib), SQL, Power BI, Excel, Git, Github.

Data Science & Analytics — Data Cleaning, Feature Engineering, KPI Analysis, ETL, EDA, Predictive Modeling.

AI & ML Foundations — Retrieval-Augmented Generation (RAG), Machine Learning, NLP, Generative AI.

Visualization & Reporting — Power Bi Dashboards, Power Query.

PROJECTS

RAG-Based PDF Q&A Application | GenAI (FastAPI, Qdrant, Gemini, Streamlit) 🔗

Problem: Enable intelligent document retrieval and question answering from unstructured PDF content.

- Built and deployed an end-to-end Retrieval-Augmented Generation (RAG) application using FastAPI, Gemini Embeddings, Qdrant Vector Database, and Streamlit for semantic PDF search and context-aware answers.
- Designed a complete document ingestion and retrieval pipeline including PDF chunking, vector embedding generation, similarity search, and cloud deployment using Render and Streamlit Community Cloud.

Pricing & Churn Prediction | Machine Learning (Scikit-learn) 🔗

Problem: Understand factors influencing pricing and customer churn in a rental marketplace

- Developed churn and pricing models achieving 87% accuracy and 0.42 RMSE, with 12% improved performance through tuning.
- Identified key retention and pricing drivers using statistical and feature importance analysis.

Text Classification & Sentiment Analysis | Python (TF-IDF, Scikit-learn, NLTK) 🔗

Problem: Classify text data for sentiment analysis and multi-category news classification.

- Built TF-IDF models on 50k+ samples, achieving ~90% (IMDb) and ~65% (news) accuracy with Logistic Regression.
- Outperformed baseline models by ~3–5% and validated using cross-validation and key metrics.

Retail Performance Analysis | SQL (CTES, Window Functions) 🔗

Problem-Analyze branch-level revenue and customer behavior to identify high-performing outlets.

- Executed advanced SQL analyses to evaluate revenue, customer patterns, and branch KPIs, identifying top 37% revenue-contributing outlets.
- Optimized query execution by 30%, improving reporting efficiency.

CERTIFICATES

AI Fluency for Students - Anthropic

Data Science Training - Internshala Trainings

Data Driven- Machine Learning Training

NLP Essentials_ Applications to Real-World Problems Training

Deep Learning Data-Intensive Training